

Cloud Computing

PaaS: Web Crawler

Sonja Biedermann Augusto Jose de Oliveira Martins
Tobias Harald Herbert

December 10, 2018

Contents

1 Project Description

2 Implementation

We decided to implement the project using Python3. Our cloud platform of choice is AWS. The only dependencies the project has are

- `requests`, for issuing HTTP requests,
- `validators`, for validating URLs, and
- `boto3`, the AWS SDK for Python.

We use DynamoDB for storage and SQS for message passing. The workers are realized as AWS Lambda functions. The master can be deployed using ElasticBeanstalk (AWS' traditional PaaS offering), which in turn uses Autoscaling, S3, Load Balancing and monitoring automatically, although this is not needed—the traffic on the master is very low.

The implementation can be found on GitHub.

2.1 Frontend

The frontend is a simple one-page website which communicates with the master node running on EC2 and displays data derived from the crawl and information about the current state of the crawler.

We display how many workers are currently working on the task. Since we use AWS Lambda, whose functions are quite short-lived, this chart sees a lot of activity with workers popping in and out of existence. The purpose of this chart is to “visualize the cloud” and see how it scales.

The result of the crawl is displayed as a network. Nodes are colored according to the host, however note that colors may be repeated. To draw the network, we use the constraint-based layout algorithm CoLa¹, which basically runs a small simulation and optimizes the

¹<https://ialab.it.monash.edu/webcola/>




Figure 1: Screenshot of the graph view

layout. The results from this algorithm are excellent, but graph drawing is hard and as a consequence we limit the graph display to 500 edges. The performance is heavily dependent on the graph structure, e.g. networks with very fat hubs (e.g. Wikipedia) have noticeably worse performance than networks with many smaller hubs. This is due to how large the neighborhoods tend to be in the hubs vicinity. It also does not help that it is running in the browser.

The graph is interactive and allows you to drag the nodes around should you not be satisfied with the layout. You can also click on a node to visit the associated webpage.

Figure ?? shows the result of a crawl on the webpage of an ongoing research project. The graph shows 4 hubs: reddit (red), arXiv (green), Wikipedia (orange), and the web presence of the University of California in Riverside (blue), which is split between two involved researchers.

The bottom displays the current number of discovered edges, which is equivalent to the amount of pages crawled, as well as the current depth. Do note that the typical web page has an extremely high fan out at depths of about 3 to 4, so the time spent at consecutive depths grows exponentially. We have thus not opted to display this as a progress bar, since the rate of progress is hard to estimate.

2.2 Backend

2.2.1 Master

The master runs a lightweight web server (using Bottle, a micro WSGI web framework) serving the frontend and offers a small REST API for querying the state of the crawl. This basically consists of starting the crawl and querying the workers and discovered edges.

The master can be deployed to an EC2 instance using ElasticBeanstalk or run locally.

2.2.2 Worker

2.2.3 Storage

The storage is realized with DynamoDB which is built for handling large amounts of reads and writes up to millions per second. As the amount of writes scale exponentially with the crawler, it has to have the ability to scale accordingly which DynamoDB provides. It works with a key:value store, is document driven and fully managed. The data we need to store is quite simple as it is at most a JSON object with three attributes stored in DynamoDB as a String. The attributes are the source link, the link found on the website of the source link and the depth or layer of the crawl. The implementation of the database consists of all function calls to write files to the database, check if an item is in the database and retrieve both a single item as well as all items in a table. These building blocks allow the master to retrieve data about all the edges required to draw the graph and it allows the workers to check if the retrieved url is already in the database and if not to write it to the database together with its source link and the depth of the crawl.

3 Setup

We have included setup scripts in `setup`. They create all resources used and also set up the needed policies.

3.1 Master

You can either run the master locally using the `master.py` contained in the project root. It serves on port 8080.

If you want to deploy it via ElasticBeanstalk, do the following at project root:

```
$ cd eb_master
$ make
```

The Makefile copies the files properly, runs the EB CLI commands to get everything ready, attaches the needed policies to your user and also the EC2 instance role which will run the application, and then launches the application in your default browser. We presume that the AWS and EB CLI are installed and set up properly.

Once the master is up and running, you can deploy the worker as described in the next section. The ordering between master and worker is not actually important.

To terminate the EC2 instance and clean up the files run `make clean` later. Please be aware that deletion will take a bit to take effect. Also, the output tells you that it is fine to Ctrl+C, but don't—while interrupting the EB CLI (who produces the output) is fine, interrupting `make` is not.

3.2 Worker(s)

3.3 Storage

The setup of the DynamoDB requires permissions of the user and the creation of the tables. Both are provided as scripts as part of the automation in the build process and are included in the Makefile.